

# RSPH HPC Cheat Sheet

Allison Codi

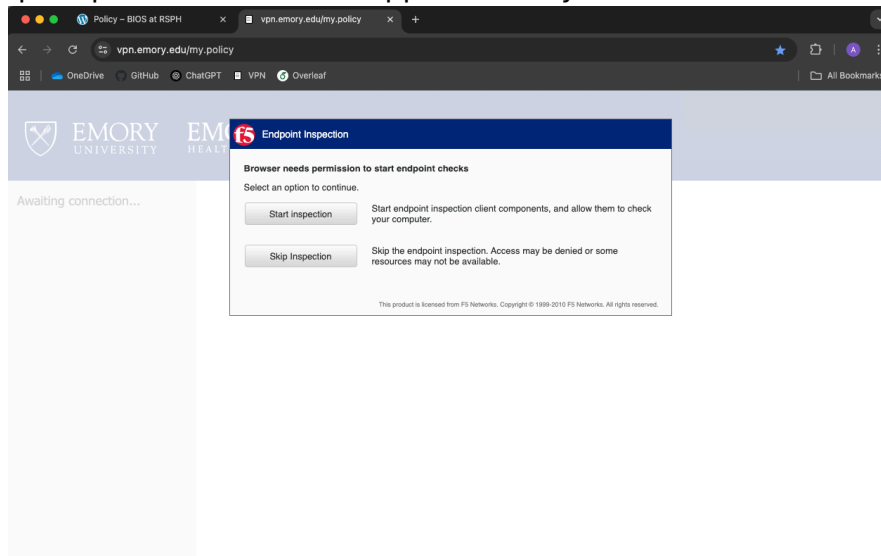
5/22/25


Official website: <https://scholarblogs.emory.edu/rsph-hpc/>

Official policy documentation: <https://github.com/RSPH-HPC/Documentation/blob/master/HPC%20Getting%20Started%20Guide%20V.2.0.pdf>

## I. Connecting to HPC


1. Connect to VPN <https://vpn.emory.edu/>
  - a. The first time you connect, it will prompt you to download the VPN client
  - b. After the client is installed, you can connect to the VPN by opening the application directly or logging in through the VPN website. Follow the prompts until the F5 VPN application says 'Connected'





EMORY

UNIVERSITY



EMORY

HEALTHCARE

VPN Remote Network Access


Emory Remote Network Access

Powered by F5 Networks

Username

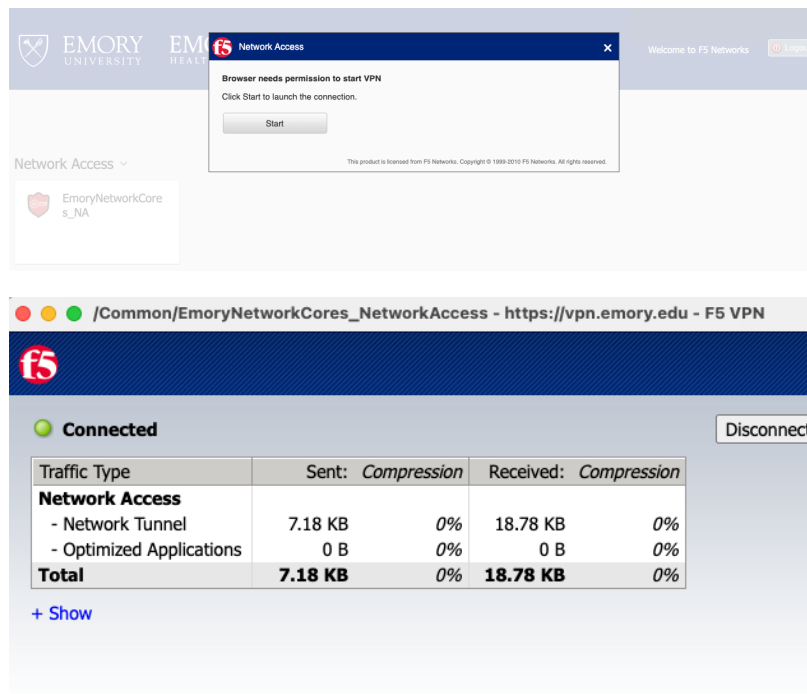
Password

Logon



ATTENTION!


Initial access automatically



2. SSH into HPC via terminal
  - a. Macbook
    - i. Open terminal
    - ii. Type `ssh <user id>@clogin01.sph.emory.edu` and hit enter
    - iii. Enter password (you will not see characters appear, but type anyway; hit enter when done)

```
Last login: Wed Feb 26 15:34:47 on ttys000
[acodi@BIO-ACODI-01M ~ % ssh acodi@clogin01.sph.emory.edu
[acodi@clogin01.sph.emory.edu's password:
```

Welcome to the



High Performance Computing (HPC) Cluster

\*\*\* AUTHORIZED USE ONLY \*\*\*

```
-->>> DO NOT RUN APPLICATIONS ON THE LOGIN NODE <<<---
        Please submit ALL computations, including small
        interactive ones, to compute nodes.
```

```
Last login: Mon Feb 24 11:24:25 2025 from 10.110.82.61
(base) [acodi@clogin01 ~]$
```

## b. Windows

- i. Download a program with a terminal emulator (ex. git, which comes with git bash) or use the already-installed OpenSSH program
- ii. Open git bash or OpenSSH program
- iii. Type `ssh <user id>@clogin01.sph.emory.edu` and hit enter
- iv. Click open, enter password when prompted (you will not see characters appear, but type anyway; hit enter when done)

## II. Navigating HPC

RSPH HPC uses a Linux operating system with SLURM for cluster management

### 1. Useful Linux commands

Command	Description
cd <dir_name>	Navigates into the folder/directory "dir_name"
ls <dir_name>	Lists all files and directories in the current directory
ls -latr	List with options. -l is list in long form (includes timestamps, permissions, size), -a is all files (including hidden files beginning with '.'), -t is sort by time with newest first, -r is reverse order while sorting. See man page for additional settings
mkdir <dir_name>	create new directory
pwd	Print path of current directory
cp <file_to_copy> <copy_to_here>	Copy file to new directory and/or rename. Include name of file in path
mv <file_to_move> <move_to_here>	Move file to another location. Include name of file in path
rm <file_name>	permanently delete/remove file
rm -r <folder_name>	permanently delete/remove folder (-r means recursive, will go through entire directory and delete)
cat <file_name>	print file
touch <file_name>	create new file with name <file_name>
vi <file_name>	open existing file in <a href="#">VIM text editor</a> or create new file/open in text editor
nano <file_name>	open existing file in <a href="#">Nano text editor</a> or create new file/open in text editor
du -sh <folder_name>	check how much storage is used by folder
Ctrl + C	kill current command
man <command_name>	open manual page for command
clear	clear screen
chmod <###> <file_or_folder_name>	change permissions on file or folder. Commonly needed for changing execution permissions of scripts. chmod 711 will give you (owner) read, write, and execute permissions while other cluster users will only be able to execute. See <a href="#">here</a> for additional details

Additional resource for learning Linux: <https://www.linuxcommand.org/index.php>

## 2. Useful SLURM specific commands

Command	Description
<code>sbatch &lt;file_name&gt;</code>	submit job to SLURM with default settings. See section V for additional details on sbatch
<code>squeue -u &lt;userid&gt;</code>	list jobs currently running from user (use your userid to list your own jobs)
<code>squeue --me</code>	list your own jobs
<code>squeue -p &lt;partition_name&gt;</code>	list jobs currently running on specific partition
<code>scancel&lt;job_id&gt;</code>	cancel job currently running
<code>scancel -u &lt;userid&gt;</code>	cancel all jobs from userid
<code>sinfo</code>	View available partitions on cluster

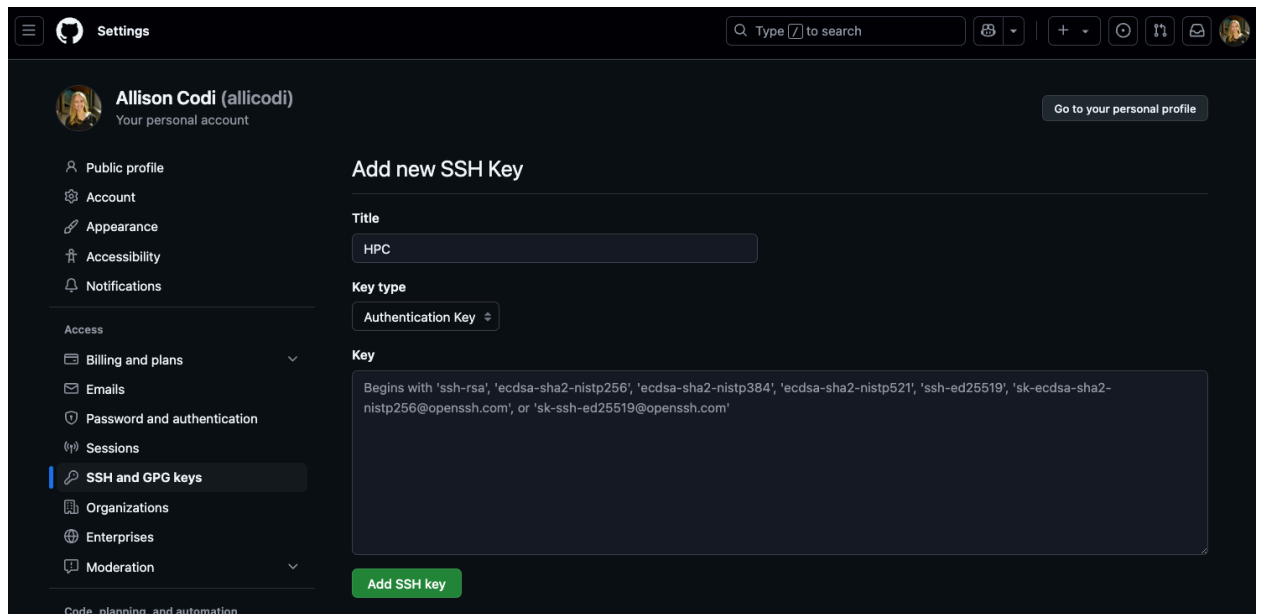
### III. Setting up ssh for GitHub

GitHub can be a helpful tool for project management and syncing files between your local computer and HPC. Upon initial setup, you must generate a ssh key and add it to your GitHub account. Detailed instructions can be found [here](#). In summary:

1. Check for existing SSH keys.
  - a. If none exist, create one using ssh-keygen (Additional instructions [here](#))
  - b. Otherwise, copy public key

```
acodi — acodi@clogin01:~/ssh — ssh acodi@clogin01.sph.emory.edu — 99x59
(base) [acodi@clogin01 ~]$ cd ~/.ssh
(base) [acodi@clogin01 .ssh]$ ls -la
total 176
drwx----- 2 acodi benkeser 4096 Mar 19 2024 .
drwxr-xr-x 26 acodi benkeser 8192 May 22 14:50 ..
-rw----- 1 acodi benkeser 568 Oct 25 2023 authorized_keys
-rw----- 1 acodi benkeser 2602 Jan 24 10:37 id_rsa
-rw-r--r-- 1 acodi benkeser 572 Jan 29 2024 id_rsa.pub
-rw-r--r-- 1 acodi benkeser 1234 Dec 16 09:24 known_hosts
(base) [acodi@clogin01 .ssh]$ cat id_rsa.pub # click enter and copy key. keep this secret! █
```

2. Log into your GitHub account. Click icon in top right corner to enter settings → SSH and GPG keys → Add new SSH key. Paste the key you copied above into the Key box, and give it an informative title.



3. Test that ssh access is working by entering `ssh -T git@github.com`. You should receive the following message:

```
(base) [acodi@clogin01 drotr_example_analysis]$ ssh -T git@github.com
Hi allicodi! You've successfully authenticated, but GitHub does not provide shell access.
(base) [acodi@clogin01 drotr_example_analysis]$ █
```

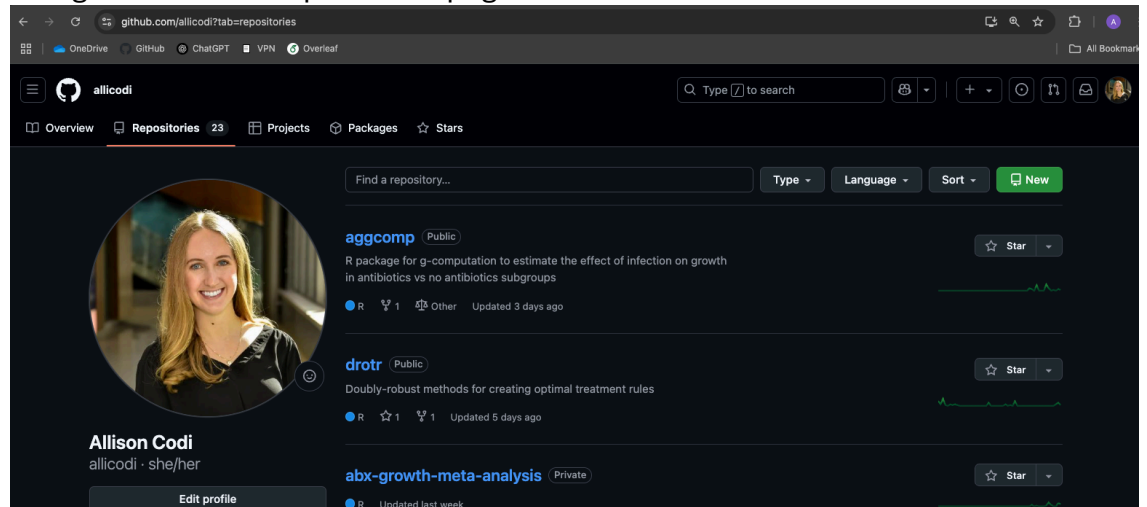
## IV. File transfer

### a. File transfer via GitHub

Useful for version control and quickly syncing code back and forth from local computer to HPC

#### 1. Set up remote repository on GitHub

##### a. Navigate to Home/Repositories page and click “New”




##### b. Enter settings for remote repository

## Create a new repository

A repository contains all project files, including the revision history. Already have a project repository elsewhere? [Import a repository](#).

Required fields are marked with an asterisk (\*).

Owner \* Repository name \*

 allicodi / example\_repo

example\_repo is available.

Great repository names are short and memorable. Need inspiration? How about [vigilant-winner](#) ?

Description (optional)

Example repo for HPC tutorial

☐ Public  
Anyone on the internet can see this repository. You choose who can commit.

☒ Private  
You choose who can see and commit to this repository.

Initialize this repository with:

☐ Add a README file  
This is where you can write a long description for your project. [Learn more about READMEs](#).

Add .gitignore


.gitignore template: None

Choose which files not to track from a list of templates. [Learn more about ignoring files](#).

Choose a license

License: None

A license tells others what they can and can't do with your code. [Learn more about licenses](#).

 You are creating a private repository in your personal account.

Create repository


allicodi / example\_repo


Type / to search


code Issues Pull requests Actions Projects Security Insights Settings

example\_repo Private

Unwatch 1 Fork 0 Star 0

**Set up GitHub Copilot**  
Use GitHub's AI pair programmer to autocomplete suggestions as you code.  
[Get started with GitHub Copilot](#)

**Add collaborators to this repository**  
Search for people using their GitHub username or email address.  
[Invite collaborators](#)

**Quick setup — if you've done this kind of thing before**  
 Set up in Desktop or **HTTPS** **SSH** `git@github.com:allicodi/example_repo.git`  
Get started by [creating a new file](#) or [uploading an existing file](#). We recommend every repository include a [README](#), [LICENSE](#), and [.gitignore](#).

**...or create a new repository on the command line**

```
echo "# example_repo" >> README.md
git init
git add README.md
git commit -m "first commit"
git branch -M main
git remote add origin git@github.com:allicodi/example_repo.git
git push -u origin main
```

**...or push an existing repository from the command line**

```
git remote add origin git@github.com:allicodi/example_repo.git
git branch -M main
git push -u origin main
```

- 
2. Create and/or link local repository to remote repository. Note this example assumes you already have set up an ssh key for GitHub to your local machine



```
example_repo — -zsh — 101x24
[acodi@BIO-ACODI-01M Documents % mkdir example_repo
[acodi@BIO-ACODI-01M Documents % cd example_repo
[acodi@BIO-ACODI-01M example_repo % git init
Initialized empty Git repository in /Users/acodi/Documents/example_repo/.git/
[acodi@BIO-ACODI-01M example_repo % touch .gitignore
[acodi@BIO-ACODI-01M example_repo % git add .gitignore
[acodi@BIO-ACODI-01M example_repo % git commit -m "init commit"
[main (root-commit) 35c23f2] init commit
 1 file changed, 0 insertions(+), 0 deletions(-)
 create mode 100644 .gitignore
[acodi@BIO-ACODI-01M example_repo % git remote add origin git@github.com:allicodi/example_repo.git
[acodi@BIO-ACODI-01M example_repo % git push -u origin main
Enumerating objects: 3, done.
Counting objects: 100% (3/3), done.
Writing objects: 100% (3/3), 213 bytes | 213.00 KiB/s, done.
Total 3 (delta 0), reused 0 (delta 0), pack-reused 0 (from 0)
To github.com:allicodi/example_repo.git
 * [new branch]      main -> main
branch 'main' set up to track 'origin/main'.
acodi@BIO-ACODI-01M example_repo %
```

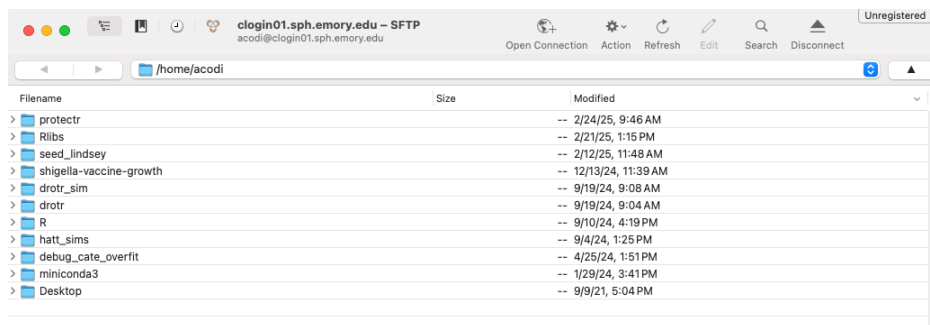
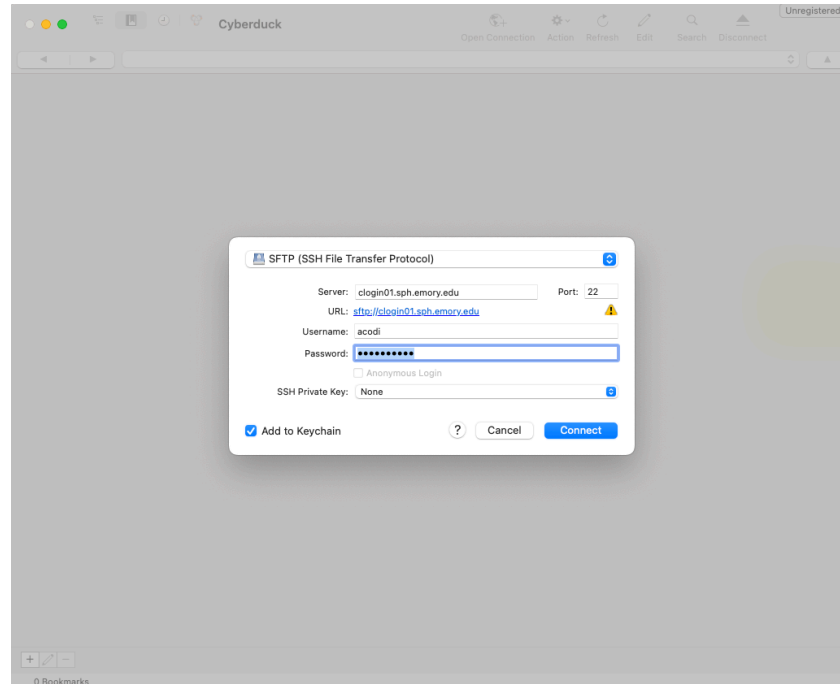
### 3. Clone remote repository on HPC

```
acodi — acodi@clogin01:~ — ssh acodi@clogin01.sph.emory.edu — 80x24
[(base) [acodi@clogin01 ~]$ git clone git@github.com:allicodi/example_repo.git
Cloning into 'example_repo'...
remote: Enumerating objects: 3, done.
remote: Counting objects: 100% (3/3), done.
remote: Total 3 (delta 0), reused 3 (delta 0), pack-reused 0 (from 0)
Receiving objects: 100% (3/3), done.
(base) [acodi@clogin01 ~]$
```

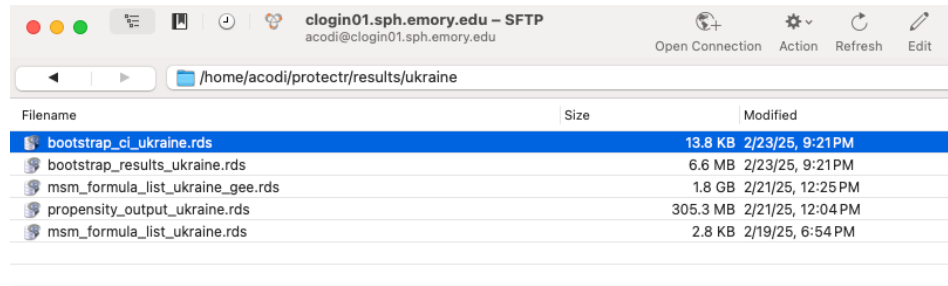
#### b. File transfer via SFTP Client

Useful for transferring large data/results files you do not want to upload to GitHub

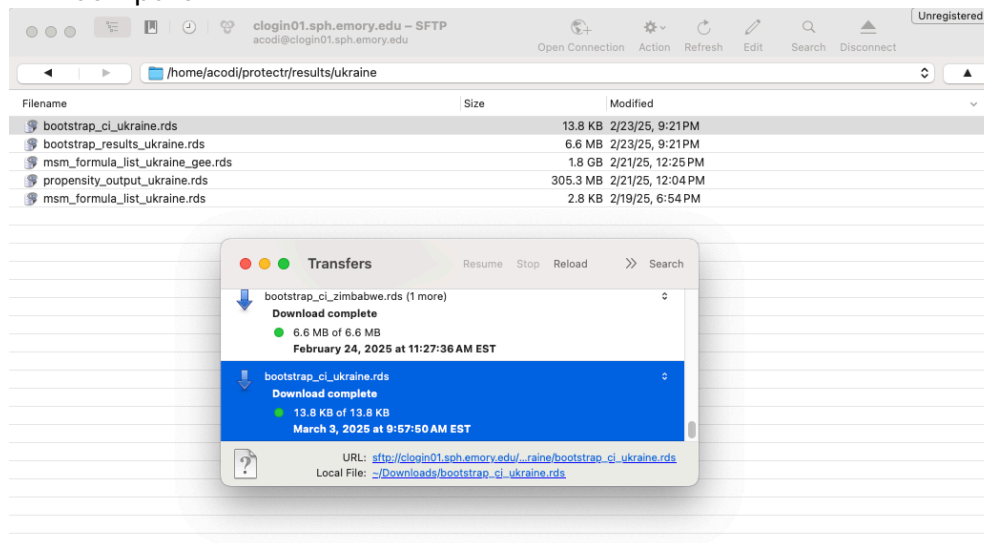
1. Download and/or open SFTP client (ex. Cyberduck)
2. Click 'Open Connection'
3. Enter ssh settings
  - a. Server = clogin01.sph.emory.edu
  - b. Port = 22
  - c. Username = <userid>
  - d. Password = password



4. Downloading files
  - a. Click into folder with file to download

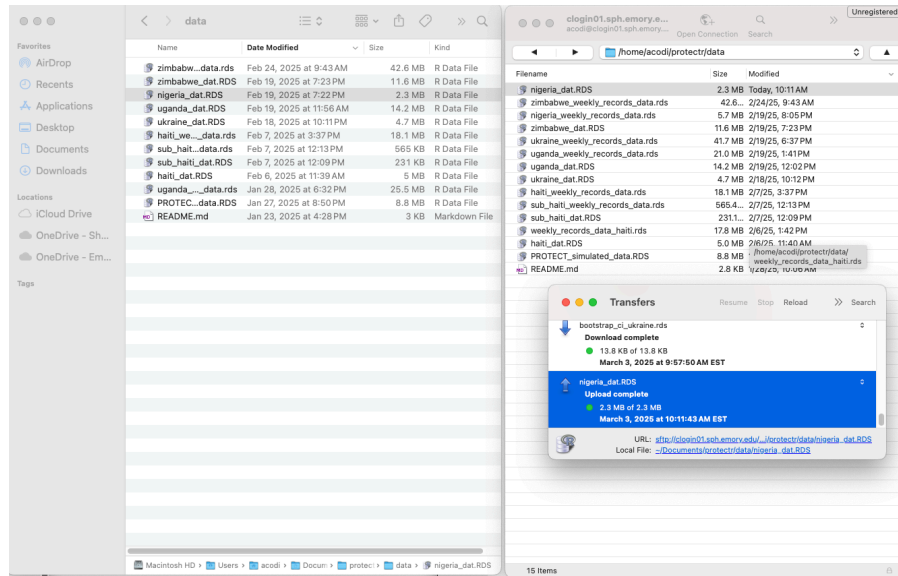


- b. Double click file to automatically download into Downloads folder on local computer



## 5. Uploading files

- a. Drag and drop file of interest from local computer into folder of interest via SFTP client



Note you can also use scp to transfer files via command line.

- Remote to local:
  - o `scp -r <your_user_name>@clogin01.sph.emory.edu:/<path_to_file>`
- Local to remote:
  - o `scp <file_name>`  
`<your_user_name>@clogin01.sph.emory.edu:/<path_to_file>`

## V. Modules

Modules refer to the software that is installed on HPC and available to use

Command	Description
<code>module available</code>	list all software installed on cluster
<code>module load &lt;module_name&gt;</code>	load a module. (ex. <code>module load R/4.4.0</code> )
<code>module list</code>	list currently loaded modules
<code>module unload &lt;module_name&gt;</code>	unload module
<code>module purge</code>	unload all modules at once

## VI. Writing sbatch scripts in bash

Bash scripts can be used to submit jobs and tailor settings using the SBATCH command

Example script 1:

```
#!/bin/bash

PARTITION=$1
SETTING=$2
NSEEDS=$3

sbatch --array=1-$NSEEDS \
    --partition=$PARTITION \
    --nodes=1 \
    --ntasks-per-node=1 \
    --cpus-per-task=32 \
    --job-name=my_job \
    --output=/projects/dbenkes/allison/my_project/scratch/${SETTING}_boot_${SLURM_ARRAY_TASK_ID}_%J.out \
    --export=SLURM_ARRAY_TASK_ID=$SLURM_ARRAY_TASK_ID,SETTING=$SETTING \
    --wrap "Rscript run_simulation.R"
```

This script runs `run_simulation.R` NSEED times on the specified partition. It will be run on a single node and use 32 CPUs on that node. Any messages/errors from the code will be saved in a scratch

file in the specified folder. The setting specified in SETTING will be available to the R script as an environment variable.

Example script 2 (from RSPH HPC documentation):

```
#!/bin/bash

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=4
#SBATCH --mem-per-cpu=1G
#SBATCH --time=01:00:00
#SBATCH --partition=week-long-cpu
#SBATCH --job-name=HelloWorld
#SBATCH --error=job.%J.err
#SBATCH --output=job.%J.out

echo "HelloWorld"

sleep 300
```

The script above will request for 1 node, 4 cores, 4Gb memory (1G per core), 1 hour run time on the partition named 'week-long-cpu'. The job name is 'HelloWorld'. This job will print 'HelloWorld' to the output file then stay idle for another 300 seconds.

Useful sbatch arguments include (but are not limited to):

Argument	Description
<b>--nodes</b>	<b>Number of nodes requested</b>
<b>--ntasks-per-node</b>	<b>Number of cores per node. For example, if --nodes=2 and --ntasks-per-node=2, then a total of 4 cores will be requested</b>
<b>--ntasks</b>	<b>Total number of cores regardless of node number</b>
<b>--mem-per-cpu</b>	<b>Memory allocated per core for the job</b>
<b>--time</b>	<b>Maximum time job can run</b>
<b>--job-name</b>	<b>Name of job, will be listed in queue</b>
<b>--partition</b>	<b>The partition to run the job on</b>
<b>--output</b>	<b>Location of stdout for the job</b>
<b>--array</b>	<b>Submit jobs over array</b>

[See here for additional arguments](#)

## VII. Interactive computing

**DO NOT RUN ANY LOCAL CODE ON THE LOGIN NODE.** If you want to run a script or test code interactively, start an interactive job on the interactive-cpu node.

```
srun --pty --partition=interactive-cpu --nodes=1 --ntasks-per-node=1 --  
mem-per-cpu=8G --time=02:00:00 bash
```

Exit interactive session with 'exit'

## VIII. Personal libraries on HPC

Some R packages are not automatically installed on HPC and may not download to the main Rlibs folder. To get around this, you can create a personal library and install R packages into it.

1. Make directory to install packages into `mkdir ~/Rlibs`
2. Set path to library at top of R script to be run on HPC using `.libPaths("~/Rlibs")`